

Assignment 5

November 1, 2022

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import statsmodels.formula.api as smf
from scipy.stats import norm
from sklearn.preprocessing import PowerTransformer
from sklearn import metrics
from scipy.stats import skew, kurtosis
import pandas_profiling as pp
```

```
[2]: df = pd.read_csv("HeartDisease.csv", index_col = "names")
```

```
[3]: df.head()
```

```
[3]:      sbp  tobacco   ldl  adiposity  famhist  typea  obesity  alcohol  age  \
names
1      160    12.00  5.73    23.11  Present    49    25.30    97.20   52
2      144     0.01  4.41    28.61   Absent    55    28.87     2.06   63
3      118     0.08  3.48    32.28  Present    52    29.14     3.81   46
4      170     7.50  6.41    38.03  Present    51    31.99    24.26   58
5      134    13.60  3.50    27.78  Present    60    25.99    57.34   49
```

```
      chd
names
1         1
2         1
3         0
4         1
5         1
```

```
[4]: df.iloc[:, :8].describe()
```

```
[4]:      count      sbp      tobacco      ldl      adiposity      typea      obesity  \
count  462.000000  462.000000  462.000000  462.000000  462.000000  462.000000
mean   138.326840   3.635649   4.740325  25.406732   53.103896   26.044113
```

std	20.496317	4.593024	2.070909	7.780699	9.817534	4.213680
min	101.000000	0.000000	0.980000	6.740000	13.000000	14.700000
25%	124.000000	0.052500	3.282500	19.775000	47.000000	22.985000
50%	134.000000	2.000000	4.340000	26.115000	53.000000	25.805000
75%	148.000000	5.500000	5.790000	31.227500	60.000000	28.497500
max	218.000000	31.200000	15.330000	42.490000	78.000000	46.580000

```

        alcohol
count  462.000000
mean    17.044394
std     24.481059
min      0.000000
25%     0.510000
50%     7.510000
75%    23.892500
max    147.190000

```

```
[5]: df.skew()
```

```

C:\Users\matth\AppData\Local\Temp\ipykernel_23600\1665899112.py:1:
FutureWarning: Dropping of nuisance columns in DataFrame reductions (with
'numeric_only=None') is deprecated; in a future version this will raise
TypeError. Select only valid columns before calling the reduction.
    df.skew()

```

```

[5]: sbp          1.180591
     tobacco      2.079210
     ldl          1.313104
     adiposity    -0.214646
     typea        -0.346438
     obesity      0.905219
     alcohol      2.312699
     age          -0.381734
     chd          0.648095
     dtype: float64

```

```
[6]: eda = pp.ProfileReport(df)
     eda.to_file("HeartDisease_Report.html")
```

```

Summarize dataset:  0%|          | 0/5 [00:00<?, ?it/s]
Generate report structure:  0%|          | 0/1 [00:00<?, ?it/s]
Render HTML:  0%|          | 0/1 [00:00<?, ?it/s]
Export report to file:  0%|          | 0/1 [00:00<?, ?it/s]

```

```
[7]: for feature in df.select_dtypes(include=np.number).iloc[:, :8]:
     mu, std = norm.fit(df[feature])
```


```


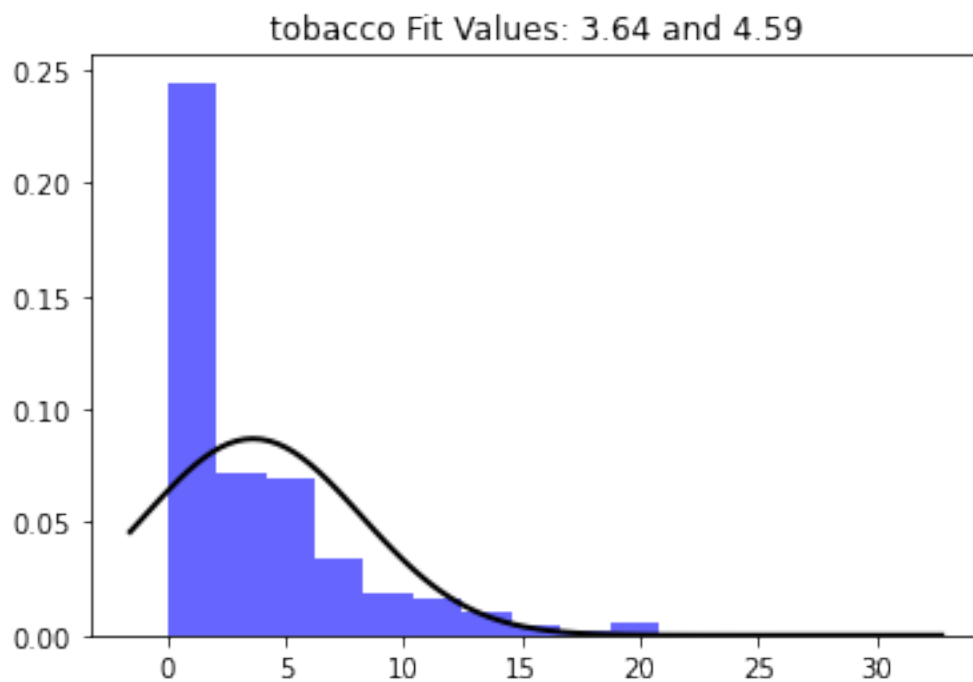
plt.hist(df[feature], bins=15, density=True, alpha=0.6, color='b')

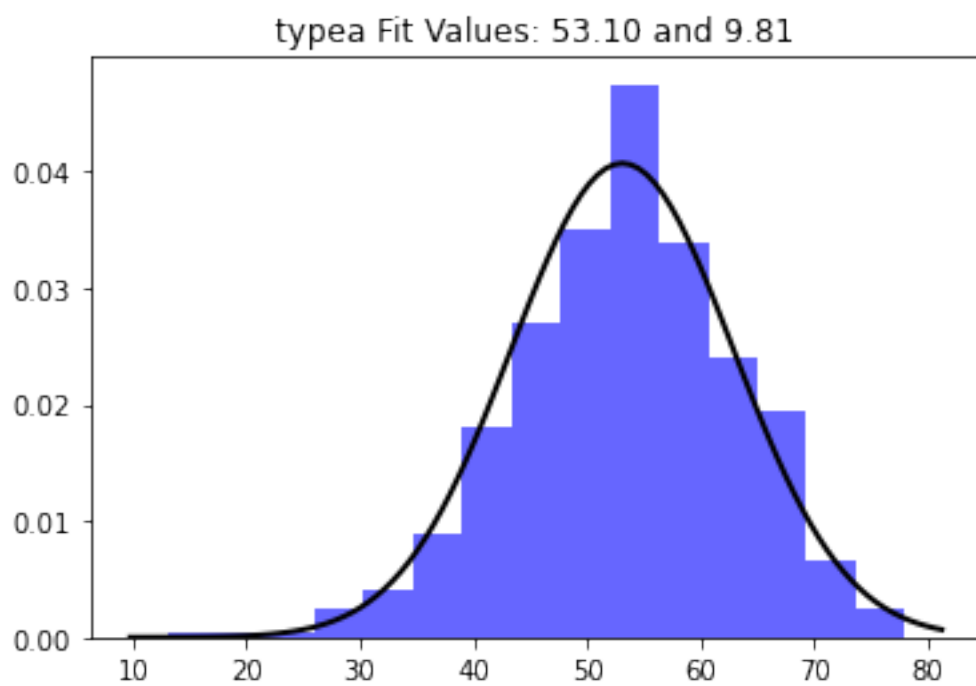

xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
p = norm.pdf(x, mu, std)

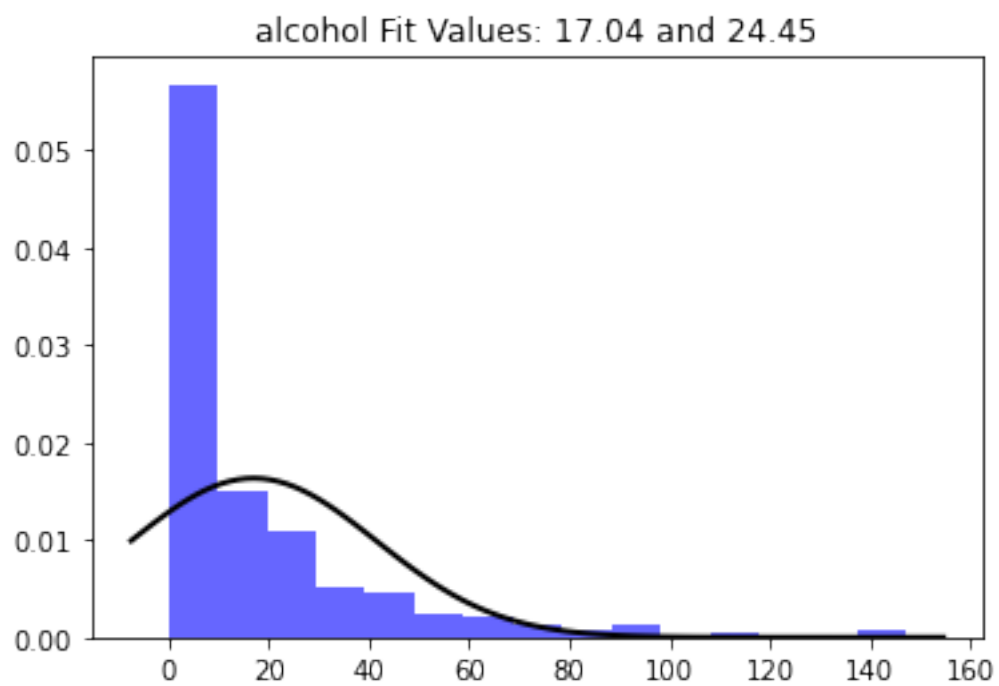

plt.plot(x, p, 'k', linewidth=2)
title = feature + " Fit Values: {:.2f} and {:.2f}".format(mu, std)
plt.title(title)
plt.savefig(feature + " histogram.jpg", dpi = 600)
plt.show()
print("Saving complete!")

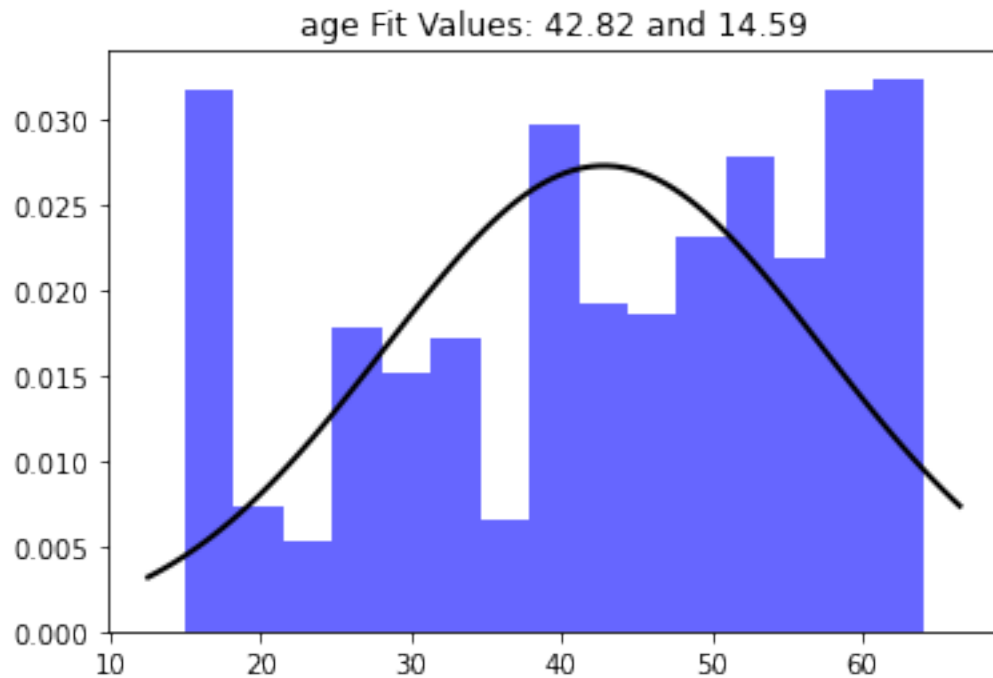
```









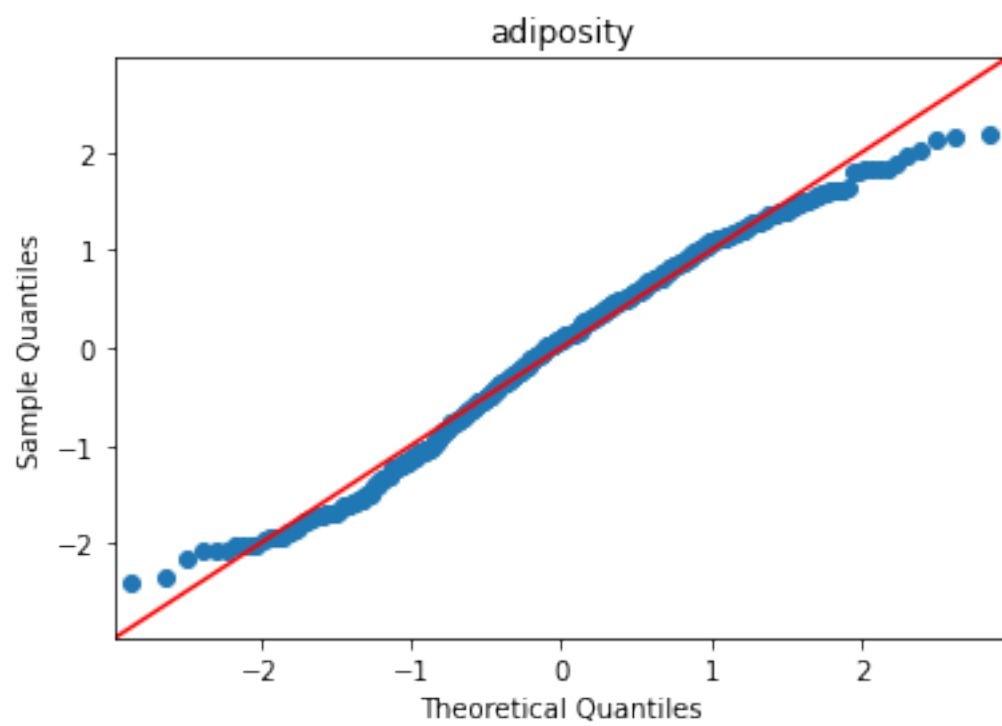
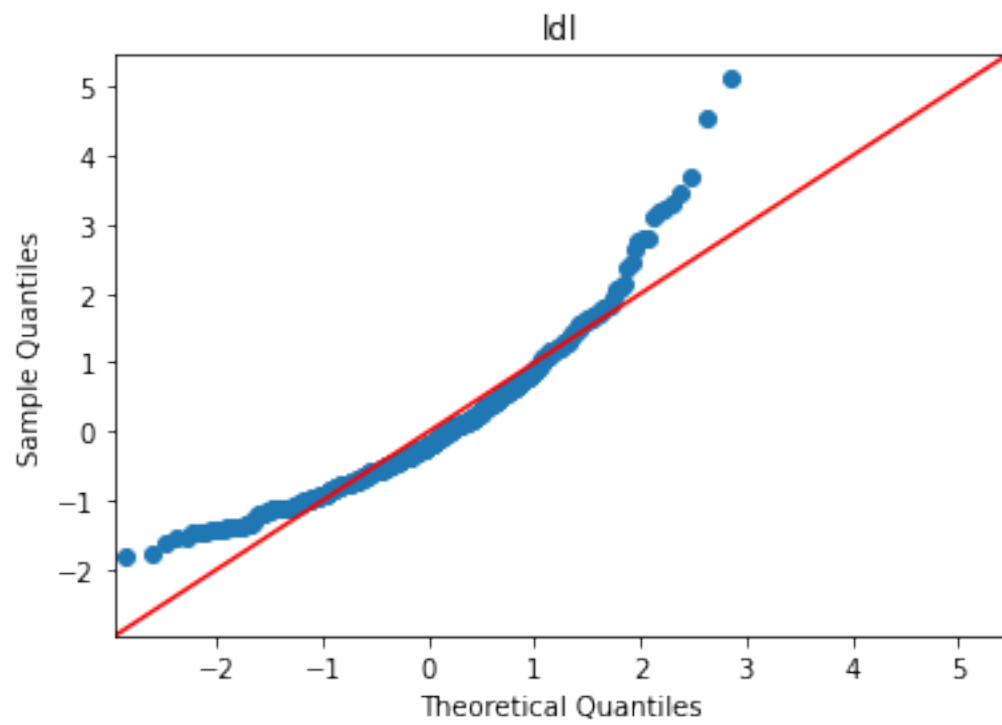


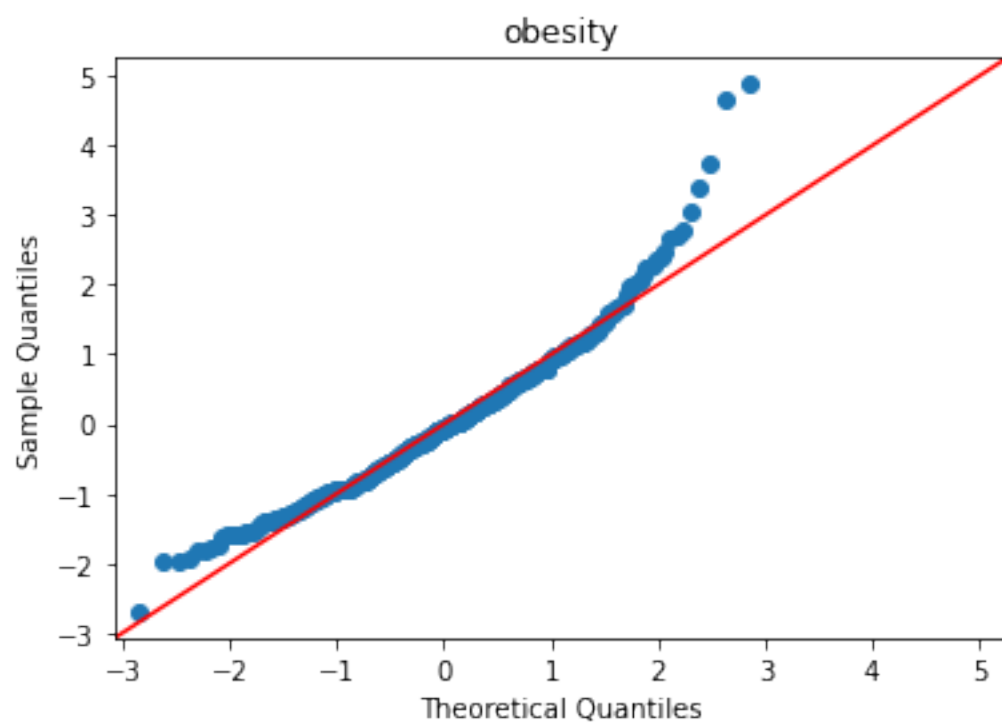





Saving complete!

```
[8]: for feature in df.select_dtypes(include=np.number).iloc[:, :8]:
      sm.qqplot(df[feature], line = "45", fit = True)
      title = feature
      plt.title(title)
      plt.savefig(feature + " qqplot.jpg", dpi = 600)
      plt.show()
```









```
[9]: df_model = pd.get_dummies(df, prefix=['famhist'])
df_model.head()
```

```
[9]:      sbp  tobacco   ldl  adiposity  typea  obesity  alcohol  age  chd  \
names
1      160    12.00  5.73    23.11    49    25.30    97.20   52   1
2      144     0.01  4.41    28.61    55    28.87     2.06   63   1
3      118     0.08  3.48    32.28    52    29.14     3.81   46   0
4      170     7.50  6.41    38.03    51    31.99    24.26   58   1
5      134    13.60  3.50    27.78    60    25.99    57.34   49   1

      famhist_Absent  famhist_Present
names
1                  0                  1
2                  1                  0
3                  0                  1
4                  0                  1
5                  0                  1
```

```
[10]: X = df_model.drop(columns = "chd")
y = df_model["chd"]
```

```
[11]: formula = "chd ~ sbp + tobacco + ldl + adiposity + typea + obesity + alcohol +
→age + famhist_Absent + famhist_Present"

log_reg = smf.glm(formula = formula, data=df_model, family=sm.families.
→Binomial()).fit()

log_reg.summary()
```

```
[11]: <class 'statsmodels.iolib.summary.Summary'>
"""
                        Generalized Linear Model Regression Results
=====
Dep. Variable:          chd      No. Observations:          462
Model:                  GLM      Df Residuals:              452
Model Family:           Binomial  Df Model:                  9
Link Function:          Logit     Scale:                   1.0000
Method:                 IRLS      Log-Likelihood:       -236.07
Date:                   Tue, 01 Nov 2022  Deviance:           472.14
Time:                   16:24:22    Pearson chi2:         452.
No. Iterations:         5          Pseudo R-squ. (CS):      0.2353
Covariance Type:        nonrobust
=====
===
                        coef      std err          z      P>|z|      [0.025
0.975]
```

```

-----
---
Intercept          -3.7920      0.871    -4.353      0.000     -5.500
-2.084
sbp                0.0065      0.006      1.135      0.256     -0.005
0.018
tobacco            0.0794      0.027      2.984      0.003      0.027
0.132
ldl                0.1739      0.060      2.915      0.004      0.057
0.291
adiposity          0.0186      0.029      0.635      0.526     -0.039
0.076
typea              0.0396      0.012      3.214      0.001      0.015
0.064
obesity            -0.0629      0.044     -1.422      0.155     -0.150
0.024
alcohol            0.0001      0.004      0.027      0.978     -0.009
0.009
age                0.0452      0.012      3.728      0.000      0.021
0.069
famhist_Absent     -2.3587      0.447     -5.279      0.000     -3.234
-1.483
famhist_Present    -1.4333      0.454     -3.159      0.002     -2.323
-0.544
=====
===
"""

```

```

[12]: df_transform = df.copy()
df_transform["sbp"] = np.power(df["sbp"], -2.0)
df_transform["tobacco"] = np.power(df["tobacco"], 0.4)
df_transform["ldl"] = np.power(df["ldl"], 0.1)
df_transform["obesity"] = np.power(df["obesity"], -0.4)
df_transform["alcohol"] = np.power(df["alcohol"], 0.4)

```

```

[13]: df_transform.head()

```

```

[13]:          sbp  tobacco      ldl  adiposity  famhist  typea  obesity  \
names
1      0.000039  2.701920  1.190736      23.11  Present    49  0.274632
2      0.000048  0.158489  1.159962      28.61   Absent    55  0.260508
3      0.000072  0.364113  1.132812      32.28  Present    52  0.259540
4      0.000035  2.238847  1.204164      38.03  Present    51  0.250031
5      0.000056  2.840636  1.133462      27.78  Present    60  0.271692

          alcohol  age  chd
names

```

1	6.238304	52	1
2	1.335202	63	1
3	1.707536	46	0
4	3.580604	58	1
5	5.051066	49	1

```
[14]: df_transform.iloc[:,8].describe()
```

```
[14]:
```

	sbp	tobacco	ldl	adiposity	typea	obesity \
count	462.000000	462.000000	462.000000	462.000000	462.000000	462.000000
mean	0.000055	1.260970	1.159087	25.406732	53.103896	0.273360
std	0.000014	0.942649	0.049116	7.780699	9.817534	0.017073
min	0.000021	0.000000	0.997982	6.740000	13.000000	0.215140
25%	0.000046	0.307415	1.126213	19.775000	47.000000	0.261865
50%	0.000056	1.319508	1.158108	26.115000	53.000000	0.272470
75%	0.000065	1.977630	1.191976	31.227500	60.000000	0.285380
max	0.000098	3.959696	1.313876	42.490000	78.000000	0.341250

	alcohol
count	462.000000
mean	2.263569
std	1.799918
min	0.000000
25%	0.763885
50%	2.239993
75%	3.558795
max	7.364637

```
[15]: df_transform.skew()
```

```
C:\Users\matth\AppData\Local\Temp\ipykernel_23600\2581206598.py:1:
FutureWarning: Dropping of nuisance columns in DataFrame reductions (with
'numeric_only=None') is deprecated; in a future version this will raise
TypeError. Select only valid columns before calling the reduction.
df_transform.skew()
```

```
[15]: sbp          0.063304
tobacco          0.134728
ldl              0.028745
adiposity       -0.214646
typea           -0.346438
obesity         -0.000273
alcohol          0.400945
age             -0.381734
chd              0.648095
dtype: float64
```

```
[16]: eda = pp.ProfileReport(df_transform)
eda.to_file("HeartDisease_transformed_Report.html")
```


```
Summarize dataset: 0%|          | 0/5 [00:00<?, ?it/s]
Generate report structure: 0%|          | 0/1 [00:00<?, ?it/s]
Render HTML: 0%|          | 0/1 [00:00<?, ?it/s]
Export report to file: 0%|          | 0/1 [00:00<?, ?it/s]
```



```
[17]: feature_list = ["sbp", "tobacco", "ldl", "obesity", "alcohol"]
for feature in feature_list:
    mu, std = norm.fit(df_transform[feature])



    plt.hist(df_transform[feature], bins=15, density=True, alpha=0.6, color='b')

    xmin, xmax = plt.xlim()
    x = np.linspace(xmin, xmax, 100)
    p = norm.pdf(x, mu, std)

    plt.plot(x, p, 'k', linewidth=2)
    title = feature + "_transformed Fit Values: {:.2f} and {:.2f}".format(mu,
→std)
    plt.title(title)
    plt.savefig(feature + "_transformed histogram.jpg", dpi = 600)
    plt.show()
print("Saving complete!")
```





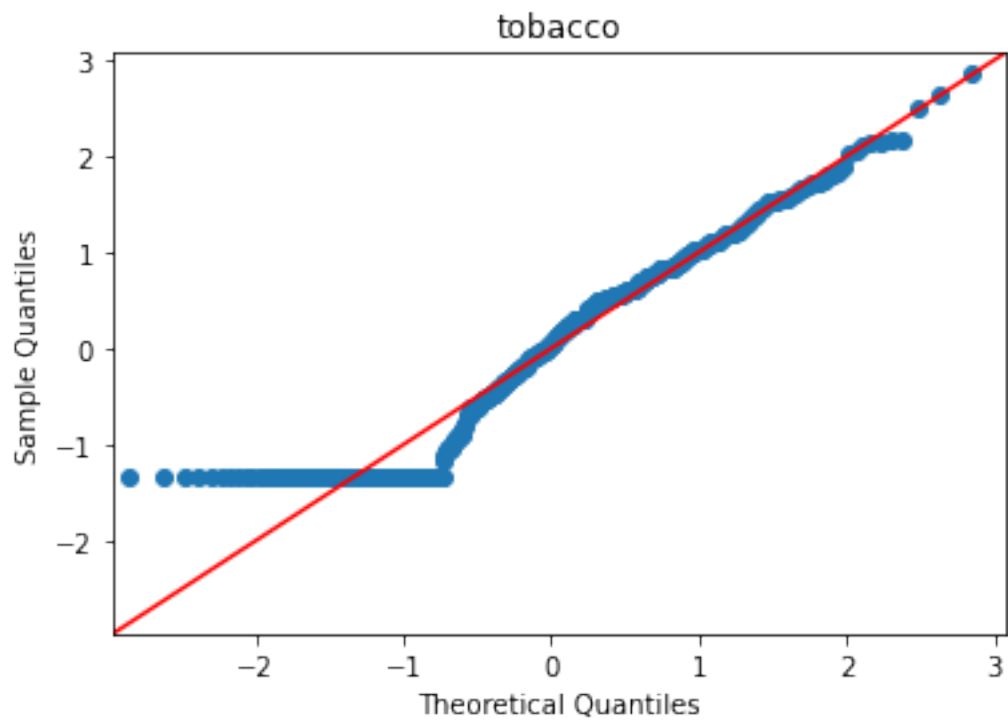


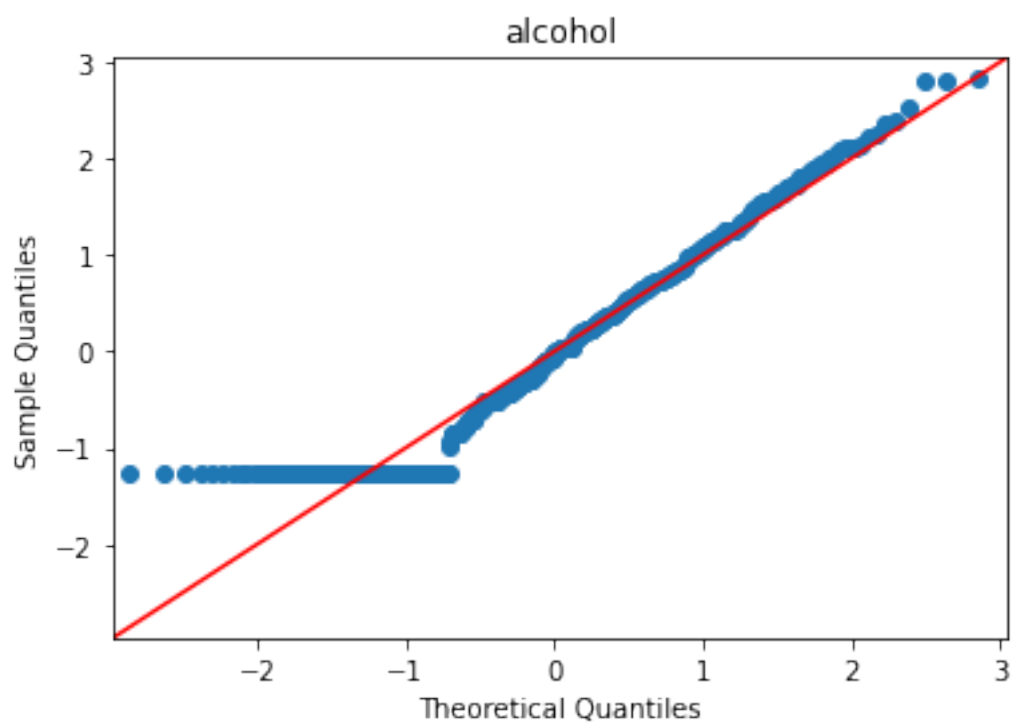



Saving complete!

```
[18]: for feature in feature_list:
      sm.qqplot(df_transform[feature], line = "45", fit = True)
      title = feature
      plt.title(title)
      plt.savefig(feature + "_transformed qqplot.jpg", dpi = 600)
      plt.show()
```







```
[19]: df_model = pd.get_dummies(df_transform, prefix=['famhist'])
df_model.head()
```

```
[19]:
```

	sbp	tobacco	ldl	adiposity	typea	obesity	alcohol	\
names								
1	0.000039	2.701920	1.190736	23.11	49	0.274632	6.238304	
2	0.000048	0.158489	1.159962	28.61	55	0.260508	1.335202	
3	0.000072	0.364113	1.132812	32.28	52	0.259540	1.707536	
4	0.000035	2.238847	1.204164	38.03	51	0.250031	3.580604	
5	0.000056	2.840636	1.133462	27.78	60	0.271692	5.051066	

	age	chd	famhist_Absent	famhist_Present
names				
1	52	1	0	1
2	63	1	1	0
3	46	0	0	1
4	58	1	0	1
5	49	1	0	1

```
[20]: X_transform = df_model.drop(columns = "chd")
```

```
[21]: formula = "chd ~ sbp + tobacco + ldl + adiposity + typea + obesity + alcohol +
→age + famhist_Absent + famhist_Present"

log_full = smf.glm(formula = formula, data=df_model, family=sm.families.
→Binomial()).fit()
log_full.summary()
```

```
[21]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        Generalized Linear Model Regression Results
=====
Dep. Variable:          chd      No. Observations:          462
Model:                  GLM      Df Residuals:              452
Model Family:           Binomial  Df Model:                  9
Link Function:          Logit    Scale:                  1.0000
Method:                 IRLS     Log-Likelihood:       -234.97
Date:                   Tue, 01 Nov 2022  Deviance:           469.95
Time:                   16:24:57    Pearson chi2:         450.
No. Iterations:         5          Pseudo R-squ. (CS):    0.2390
Covariance Type:        nonrobust
=====
===
                        coef      std err          z      P>|z|      [0.025
0.975]
-----
---
```

Intercept	-13.9631	3.352	-4.165	0.000	-20.534
-7.392					
sbp	-4800.4981	8559.784	-0.561	0.575	-2.16e+04
1.2e+04					
tobacco	0.4414	0.141	3.124	0.002	0.164
0.718					
ldl	7.6975	2.717	2.833	0.005	2.372
13.023					
adiposity	0.0345	0.030	1.145	0.252	-0.025
0.094					
typea	0.0392	0.012	3.193	0.001	0.015
0.063					
obesity	22.4078	11.475	1.953	0.051	-0.082
44.898					
alcohol	0.0012	0.065	0.019	0.985	-0.126
0.128					
age	0.0422	0.012	3.397	0.001	0.018
0.067					
famhist_Absent	-7.4285	1.681	-4.420	0.000	-10.722
-4.135					
famhist_Present	-6.5346	1.680	-3.891	0.000	-9.826
-3.243					

=====
===
"""

```
[22]: formula = "chd ~ tobacco + ldl + typea + age + famhist_Absent + famhist_Present"

log_reduced = smf.glm(formula = formula, data=df_model, family=sm.families.
    ↪Binomial()).fit()

log_reduced.summary()
```

```
[22]: <class 'statsmodels.iolib.summary.Summary'>
      """

              Generalized Linear Model Regression Results
=====
Dep. Variable:          chd      No. Observations:          462
Model:                GLM      Df Residuals:              456
Model Family:          Binomial  Df Model:                  5
Link Function:          Logit    Scale:                  1.0000
Method:                IRLS     Log-Likelihood:       -237.21
Date:                  Tue, 01 Nov 2022    Deviance:            474.43
Time:                  16:24:57    Pearson chi2:         464.
No. Iterations:         5        Pseudo R-squ. (CS):    0.2316
Covariance Type:        nonrobust
=====
```

```

===
                coef      std err          z      P>|z|      [0.025
0.975]
-----
---
Intercept      -8.9966      1.962      -4.585      0.000     -12.842
-5.151
tobacco         0.4430      0.135       3.283      0.001       0.179
0.708
ldl             6.9682      2.465       2.826      0.005       2.136
11.800
typea          0.0366      0.012       3.010      0.003       0.013
0.060
age            0.0486      0.010       4.657      0.000       0.028
0.069
famhist_Absent  -4.9365      0.981      -5.030      0.000     -6.860
-3.013
famhist_Present -4.0601      0.994      -4.086      0.000     -6.007
-2.113
=====
===
"""

```

```
[23]: X_transform_reduce = df_model.drop(columns = ["chd", "sbp", "adiposity",
→"obesity", "alcohol"])
```

```
[24]: y1 = log_reg.predict(X)
y2 = log_full.predict(X_transform)
y3 = log_reduced.predict(X_transform_reduce)
```

```
[25]: fpr1, tpr1, _ = metrics.roc_curve(y, y1)
fpr2, tpr2, _ = metrics.roc_curve(y, y2)
fpr3, tpr3, _ = metrics.roc_curve(y, y3)
auc1 = metrics.auc(fpr1, tpr1)
auc2 = metrics.auc(fpr2, tpr2)
auc3 = metrics.auc(fpr3, tpr3)
```

```
[26]: fig, axs = plt.subplots(3, figsize=(10, 30))
fig.suptitle('ALL ROC Plots of Model')
axs[0].plot(fpr1,tpr1,label="auc="+str(auc1))
axs[0].set_title('Full No Transform')

axs[1].plot(fpr2,tpr2,label="auc="+str(auc2))
axs[1].set_title('Full Transform')

axs[2].plot(fpr3,tpr3,label="auc="+str(auc3))
axs[2].set_title('Reduced Transform')
```

```
for ax in axs.flat:
    ax.set(xlabel='False Positive Rate', ylabel='True Positive Rate')
    ax.legend(loc=4)
plt.savefig("AUC Curves.jpg", dpi = 600)
plt.show()
```


ALL ROC Plots of Model

